CPSC 340: Machine Learning and Data Mining

Fundamentals of Learning

Fall 2017

Admin

- Assignment 0 is due Friday: you should be almost done.
- Waiting list people: you should be registered.
  - You may be e-mailed about prereqs, follow instructions to stay registered.
- Tutorials:
  - If sections are full, sign up for T1Z (doesn't conflict with anything).
- Important webpages:
  - www.cs.ubc.ca/~schmidtm/Courses/340-F17
  - www.piazza.com/ubc.ca/winterterm12017/cpsc340/home
  - https://www.cs.ubc.ca/getacct
- Auditing: message me on Piazza if you want to audit.
  - Bring your forms to me in class or instructor office hours.

Motivation: Determine Home City

- We are given data from 248 homes.
- For each home/object, we have these features:
  - Elevation.
  - Year.
  - Bathrooms
  - Bedrooms.
  - Price.
  - Square feet.
- Goal is to build a program that predicts SF or NY.

This example and images of it come from:
http://www.r2d3.us/visual-intro-to-machine-learning-part-1

Plotting Elevation

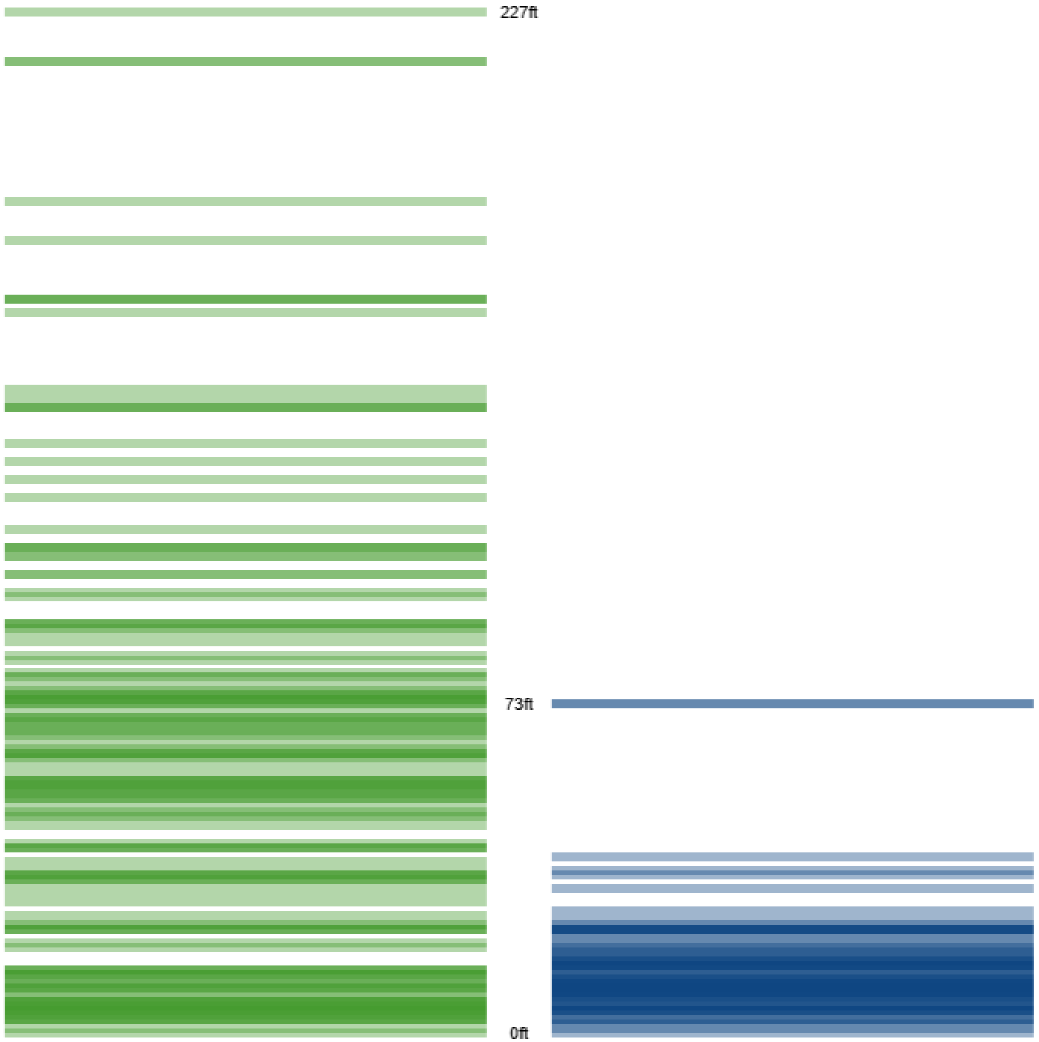

Simple Decision Stump

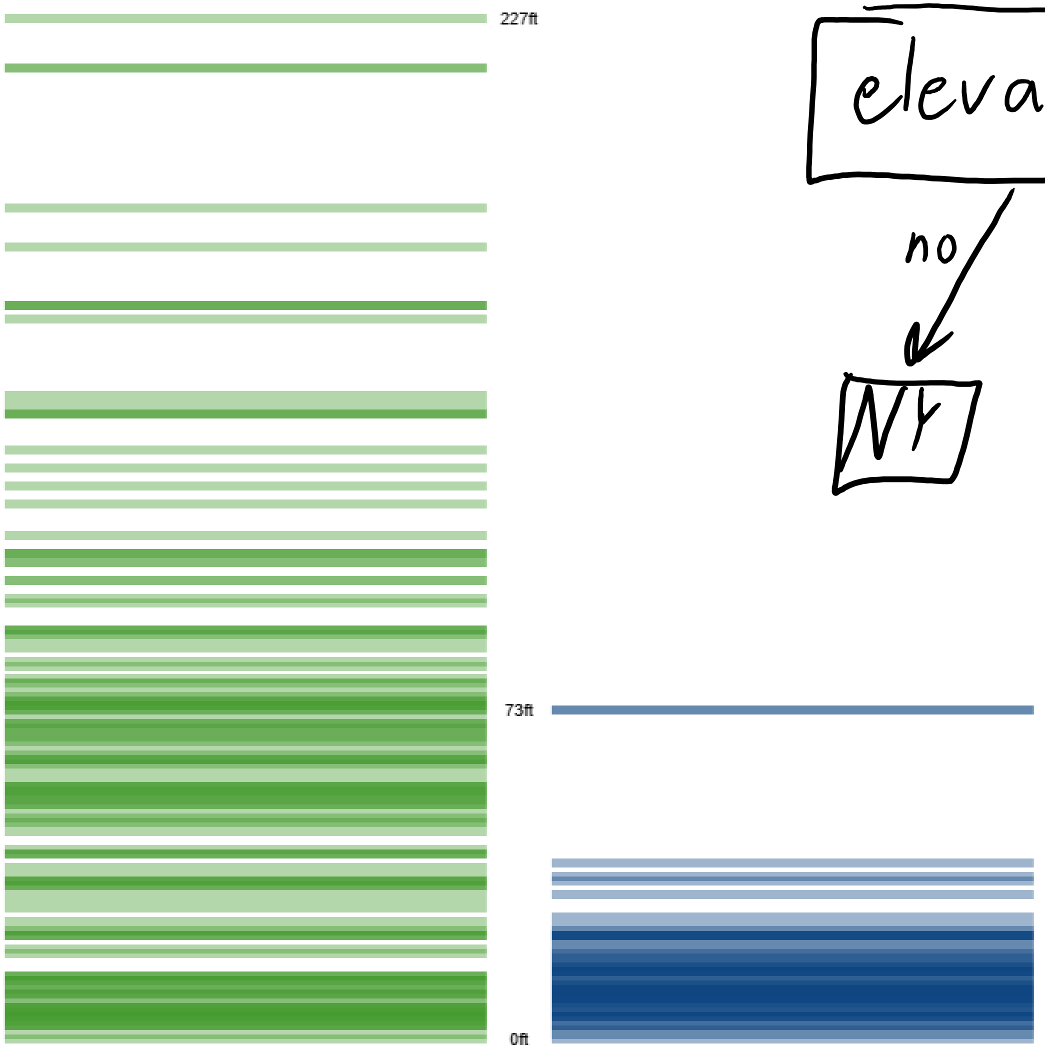

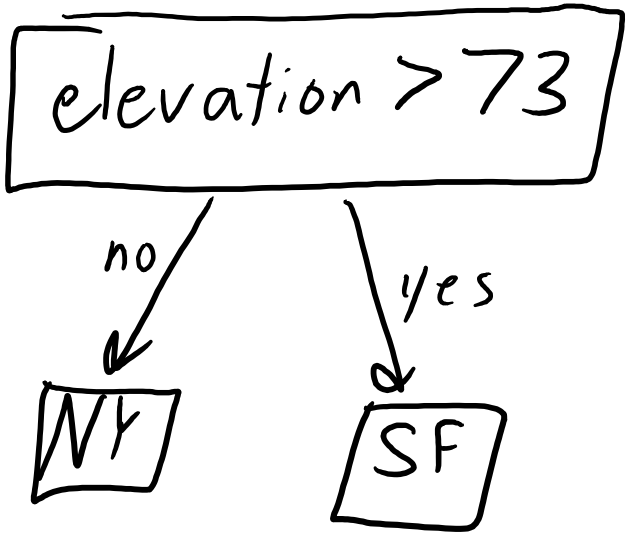

Scatterplot Array

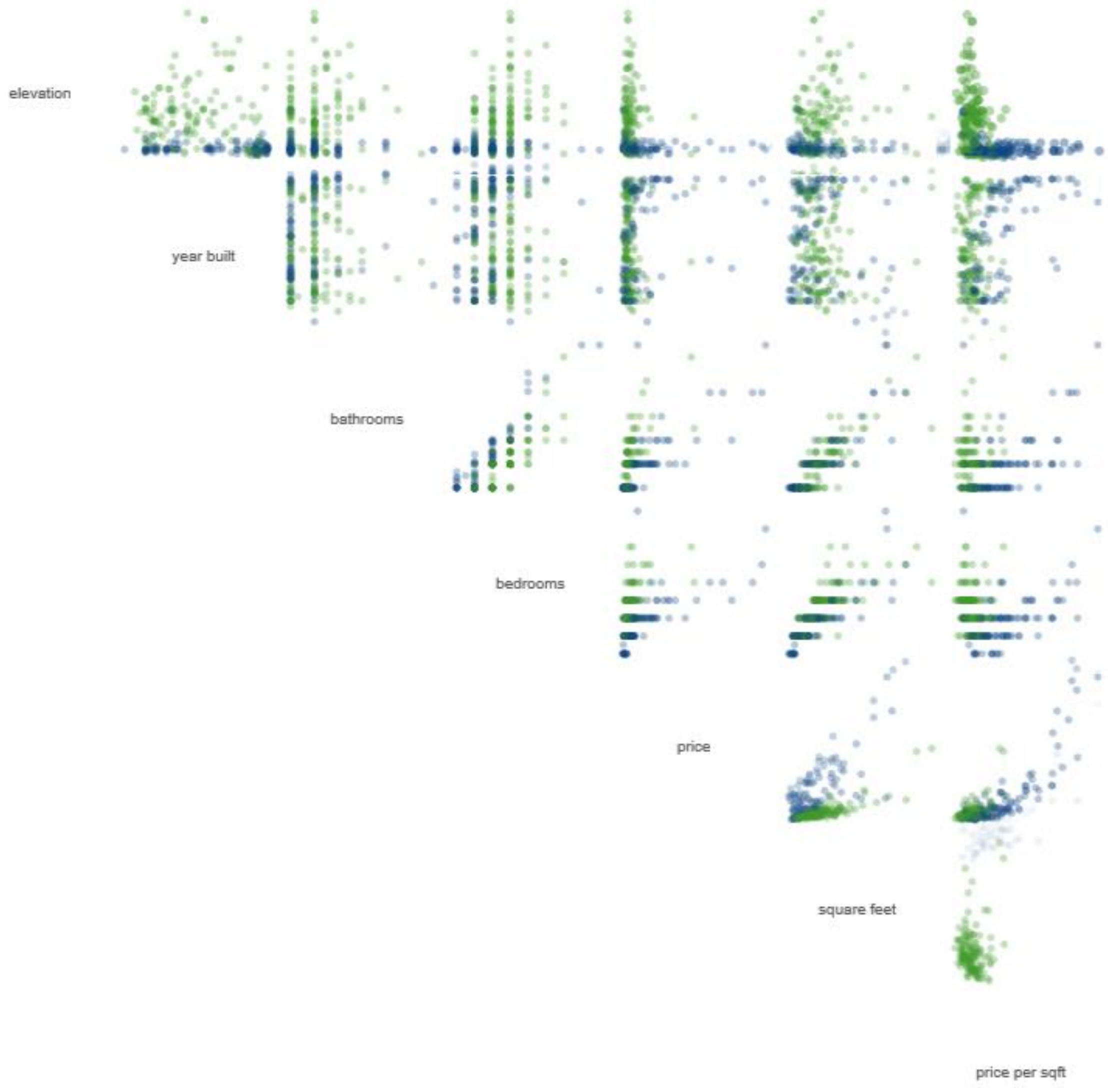

Scatterplot Array

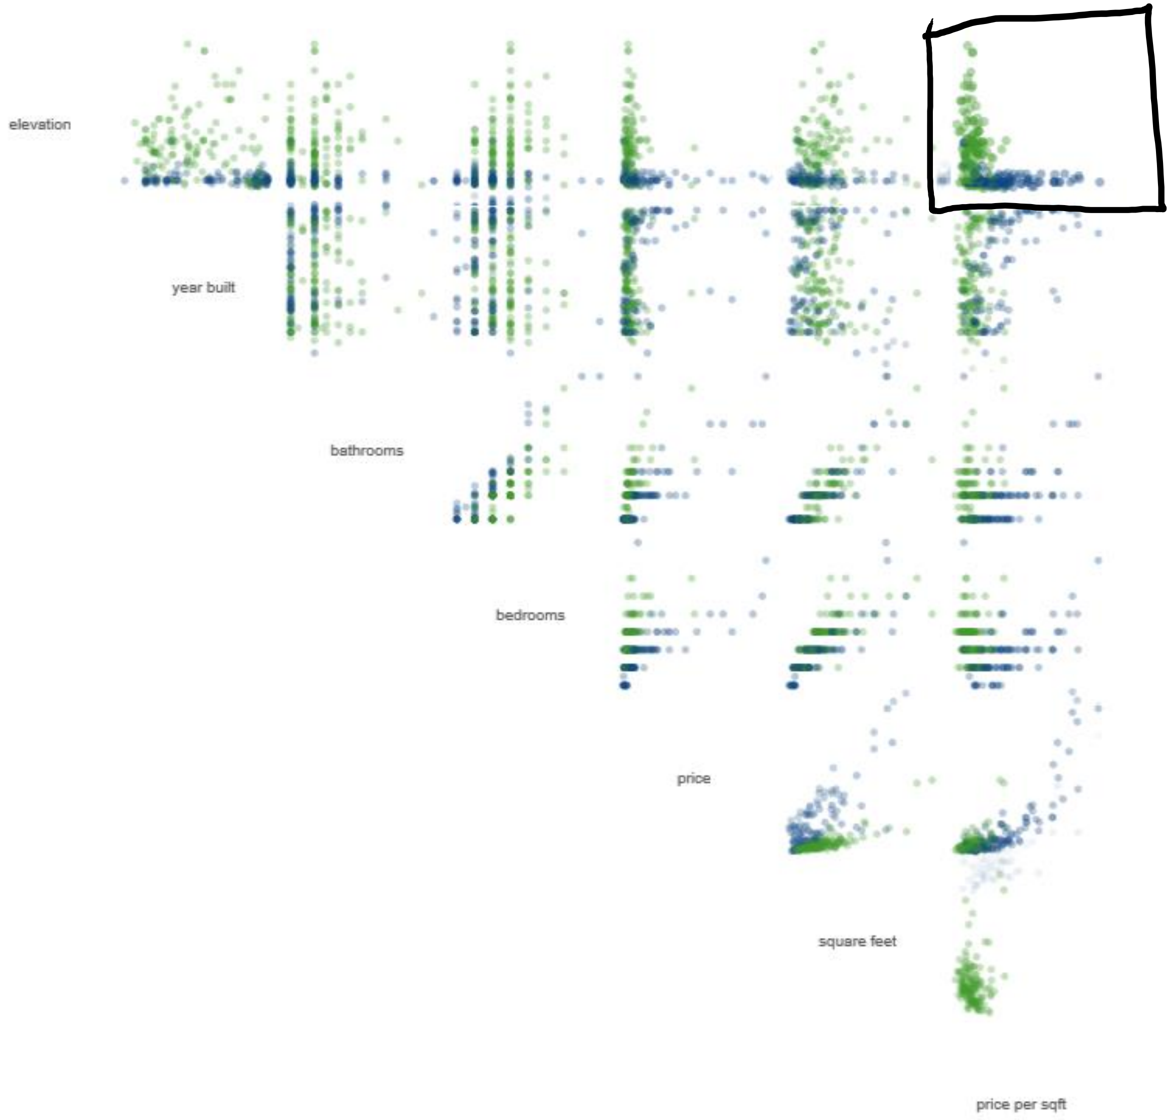

Plotting Elevation and Price/SqFt

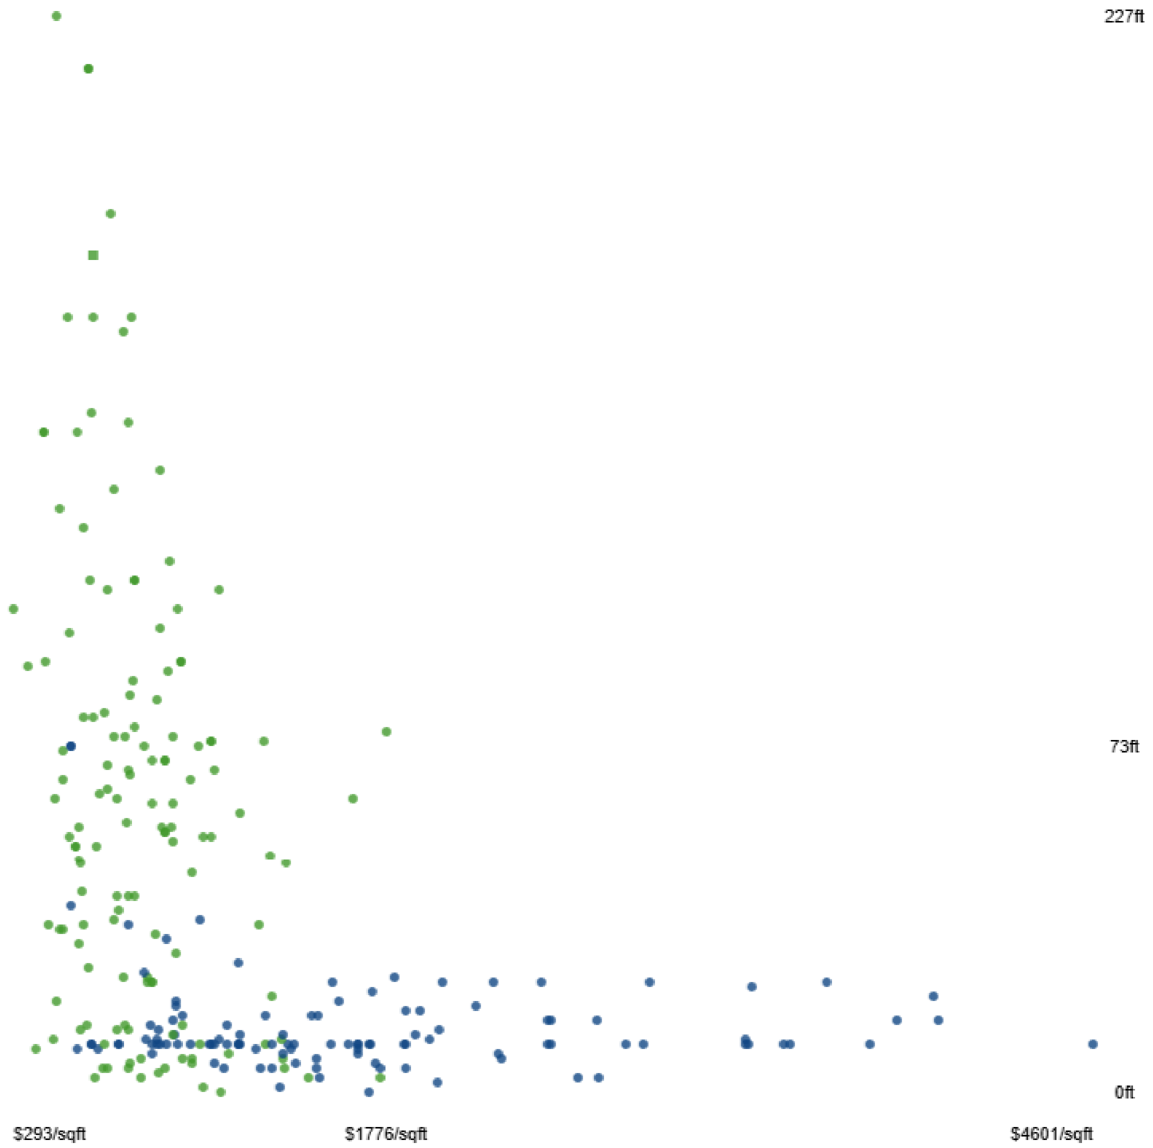

Simple Decision Tree Classification

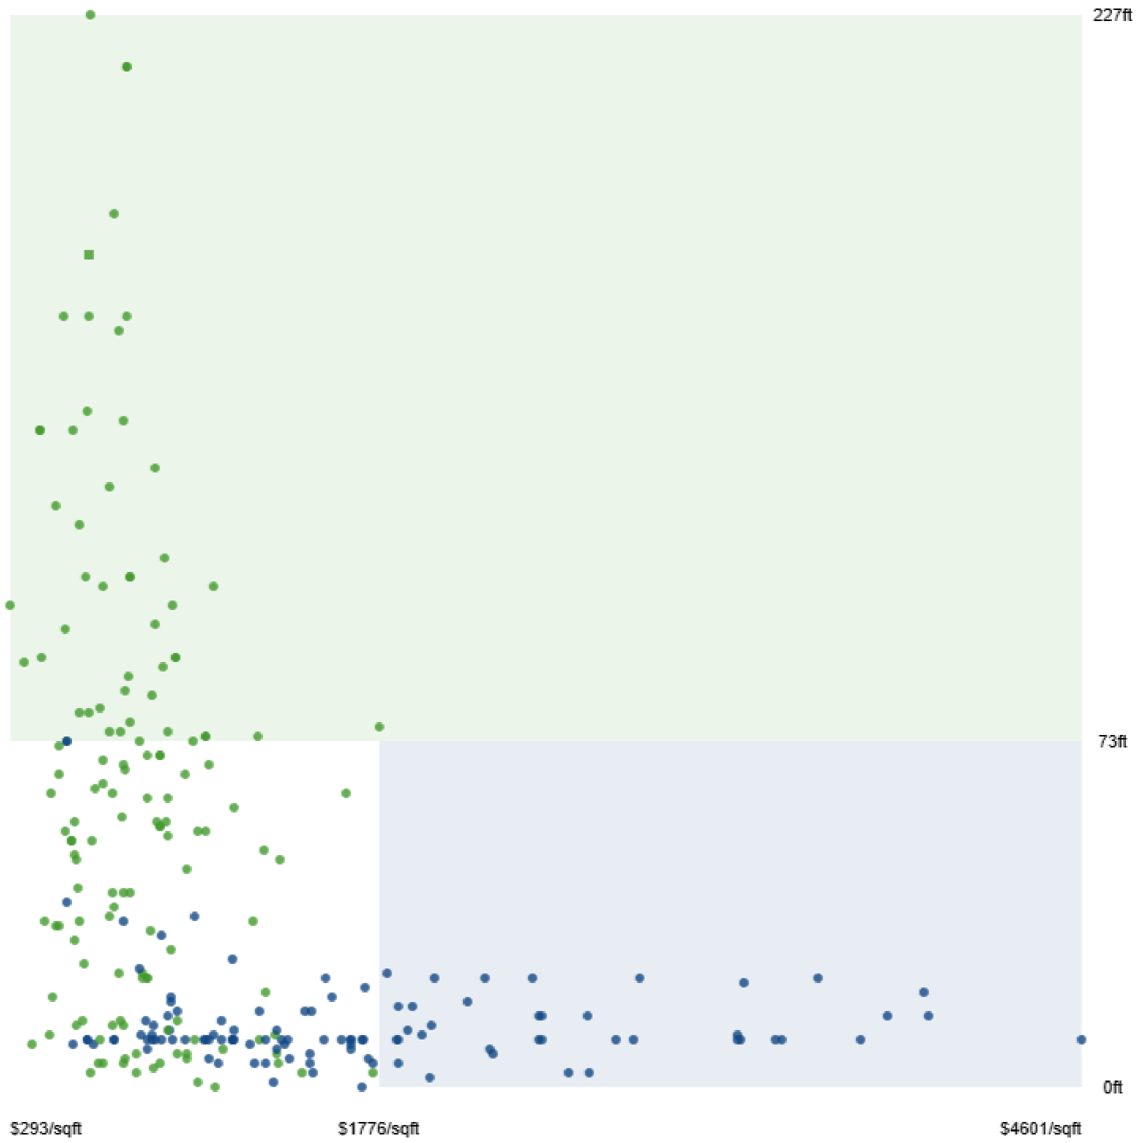

Simple Decision Tree Classification

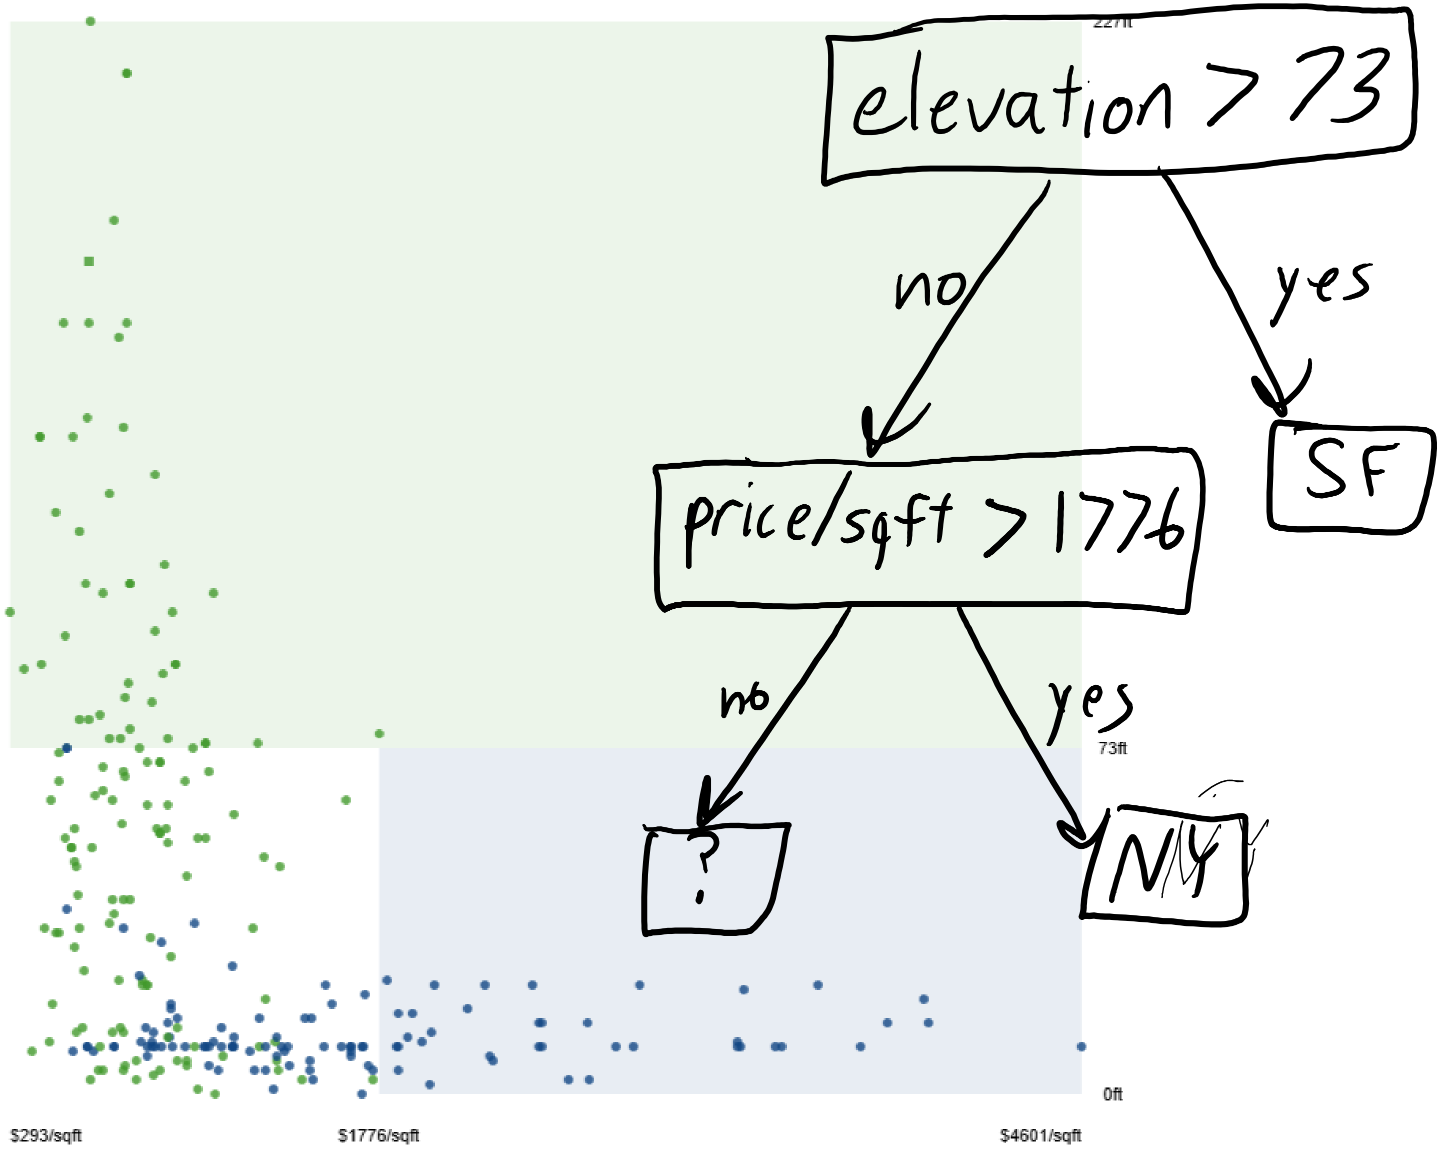

How does the depth affect accuracy?

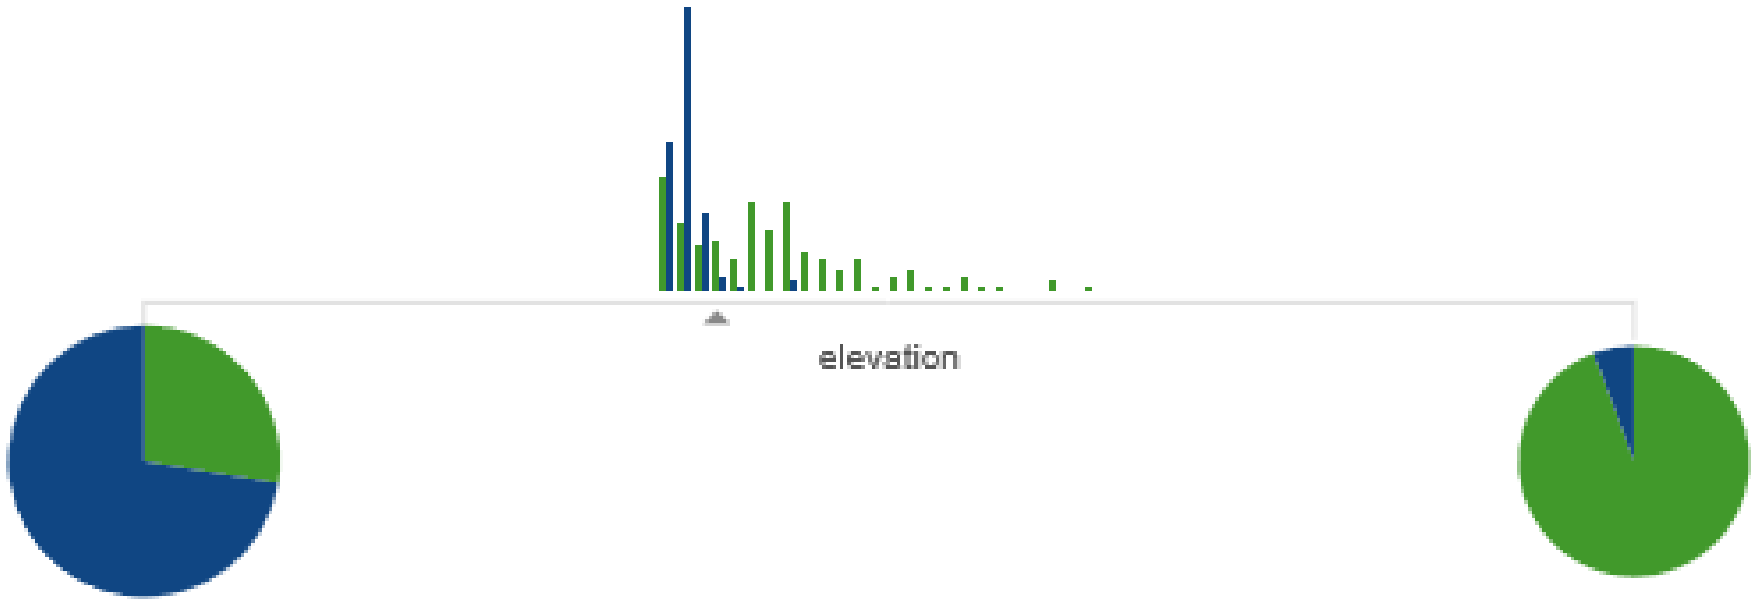

This is a good start (> 75% accuracy).

How does the depth affect accuracy?

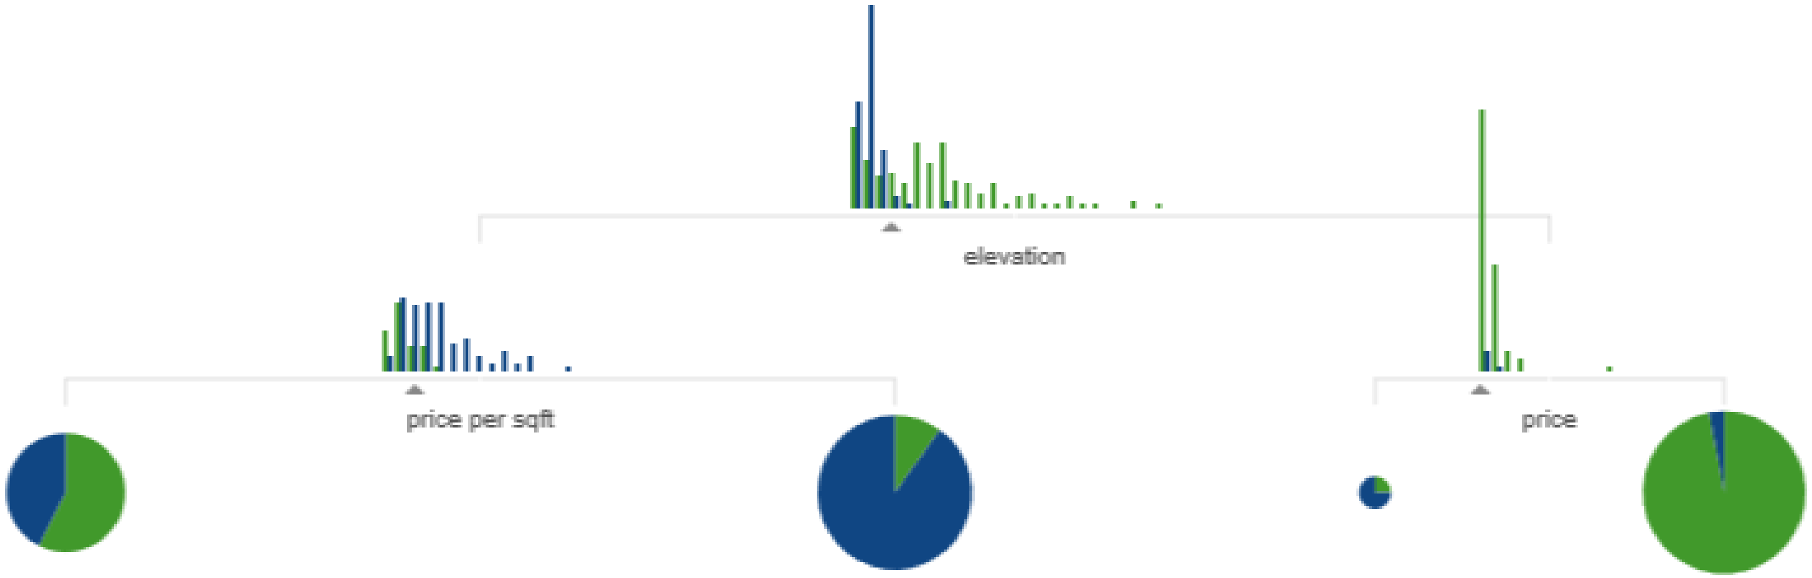

Start splitting the data recursively…

How does the depth affect accuracy?

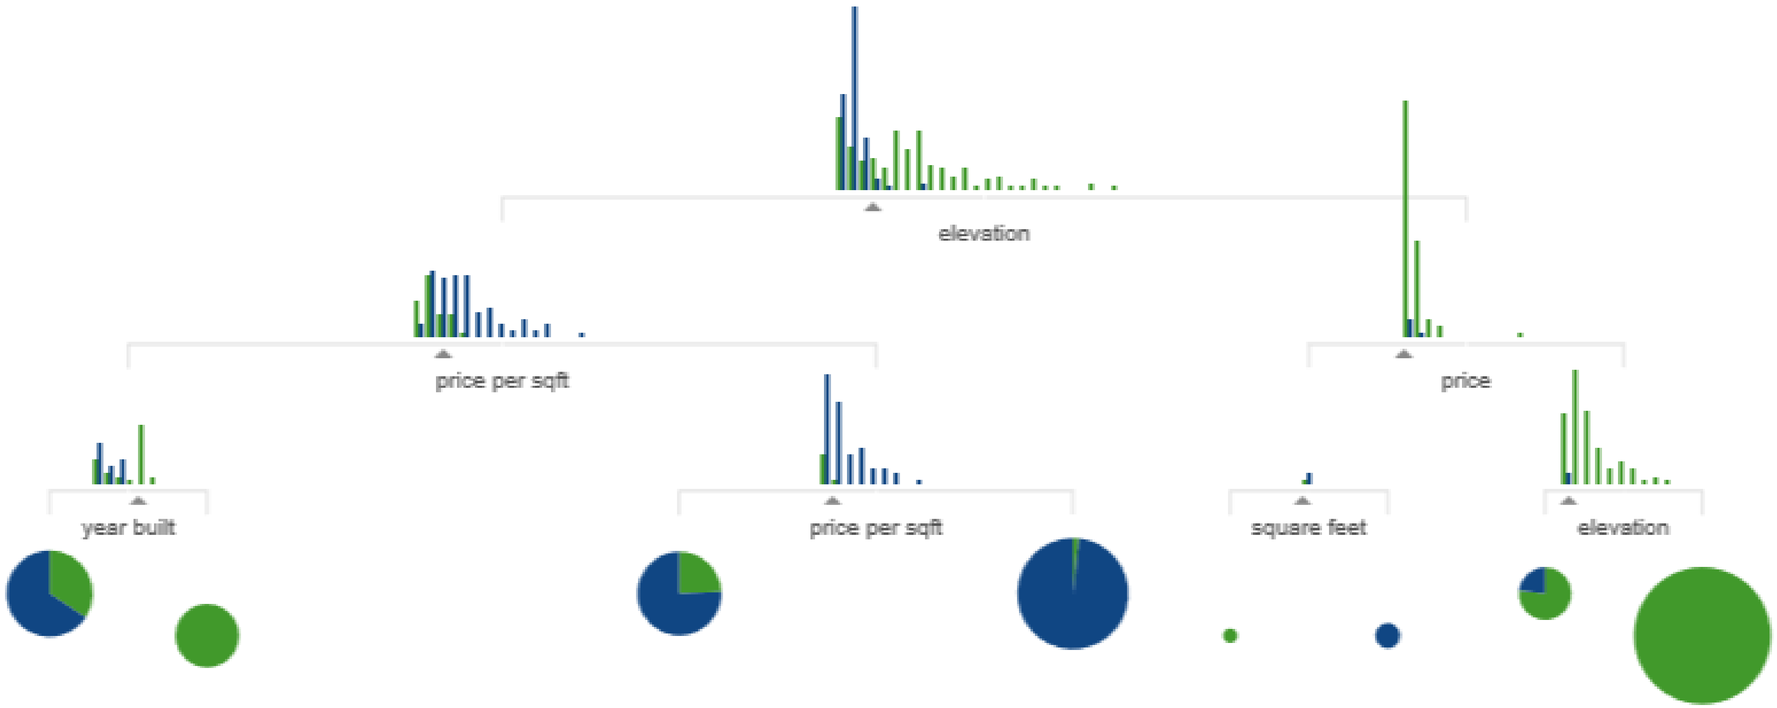

Accuracy keeps increasing as we add depth.

How does the depth affect accuracy?

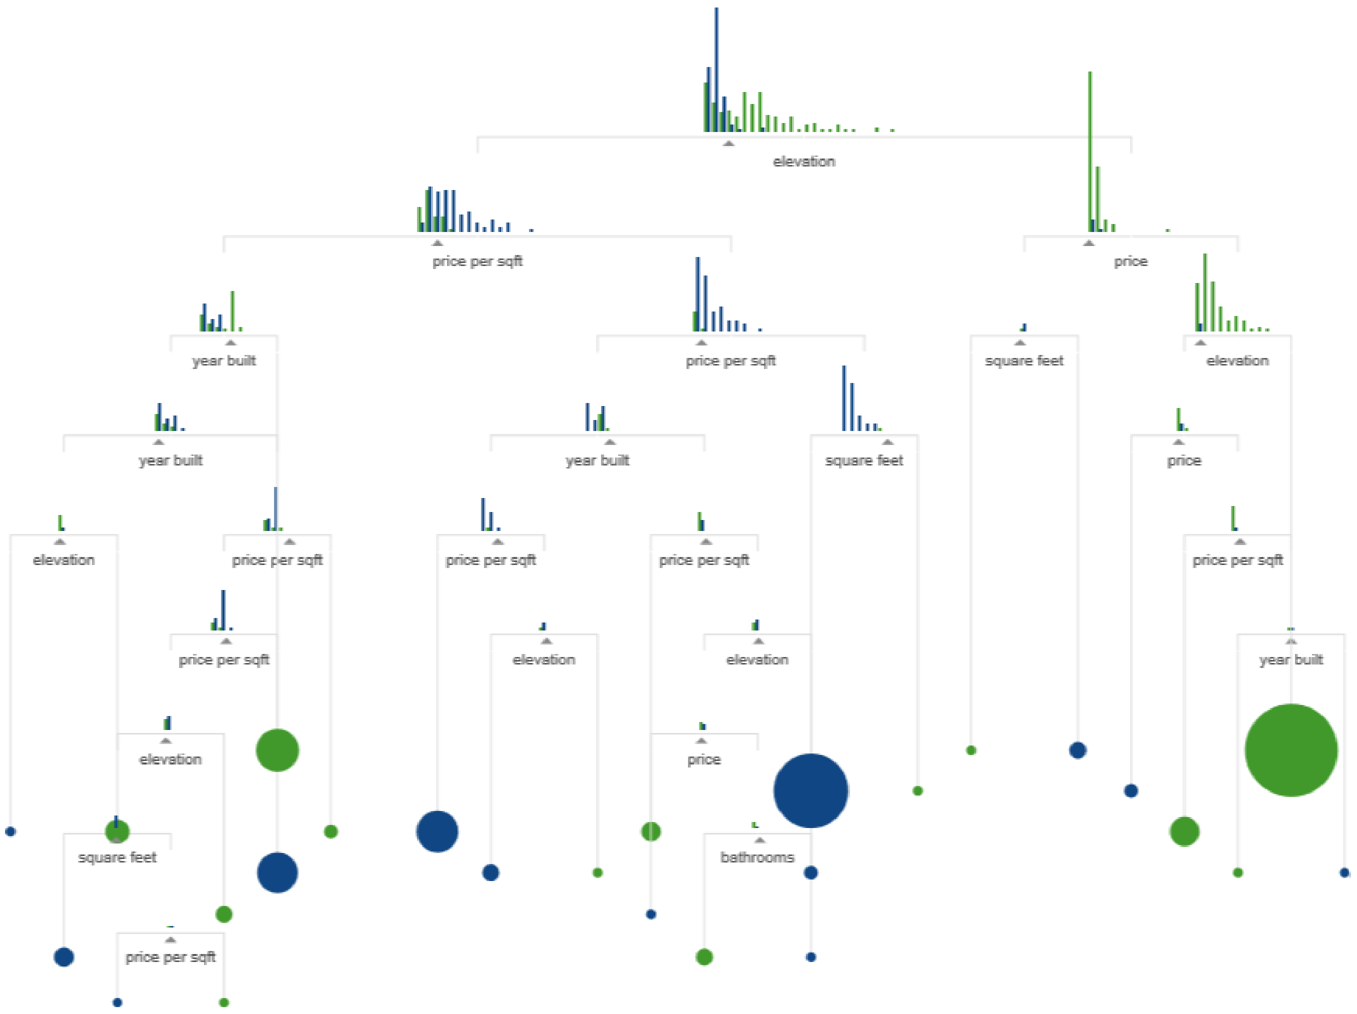

Eventually, we can perfectly classify all of our data.

Training vs. Testing Error

- With this decision tree, ‘training accuracy’ is 1.
  - It perfectly labels the data we used to make the tree.
- We are now given features for 217 new homes.
- What is the ‘testing accuracy’ on the new data?
  - How does it do on data not used to make the tree?

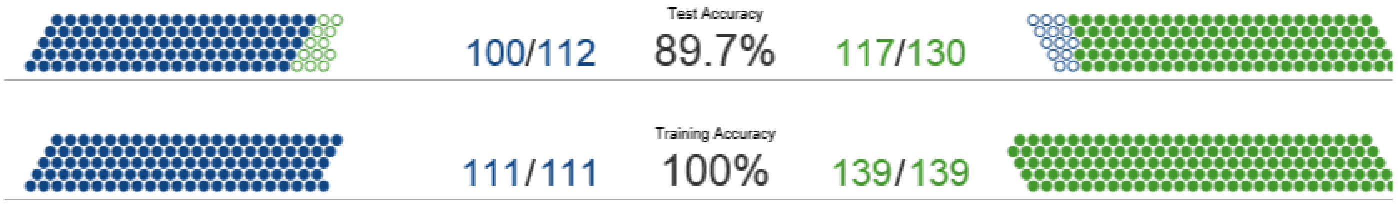

- Overfitting: lower accuracy on new data.
  - Our rules got too specific to our exact training dataset.

Supervised Learning Notation

- We are given training data where we know labels:

X =

Egg	Milk	Fish	Wheat	Shellfish	Peanuts	…
0	0.7	0	0.3	0	0	
0.3	0.7	0	0.6	0	0.01	
0	0	0	0.8	0	0	
0.3	0.7	1.2	0	0.10	0.01	
0.3	0	1.2	0.3	0.10	0.01	

y =

Sick?
1
1
0
1
1

- But there is also testing data we want to label:

Xtest =

Egg	Milk	Fish	Wheat	Shellfish	Peanuts	…
0.5	0	1	0.6	2	1	
0	0.7	0	1	0	0	
3	1	0	0.5	0	0	

ytest =

Sick?
?
?
?

Supervised Learning Notation

- Typical supervised learning steps:
  1. Build model based on training data X and y.
  2. Model makes predictions 'yhat' on test data Xtest.

$\longrightarrow \hat{y}$

- Instead of training error, consider test error:
  - Is yhat similar to true unseen ytest?

Goal of Machine Learning

- In machine learning:
  - What we care about is the test error!
- Midterm analogy:
  - The training error is the practice midterm.
  - The test error is the actual midterm.
  - Goal: do well on actual midterm, not the practice one.
- Memorization vs learning:
  - Can do well on training data by memorizing it.
  - You've only learned if you can do well in new situations.

Golden Rule of Machine Learning

- Even though what we care about is test error:
  - THE TEST DATA CANNOT INFLUENCE THE TRAINING PHASE IN ANY WAY.
- We're measuring test error to see how well we do on new data:
  - If used during training, doesn't measure this.
  - You can start to overfit if you use it during training.
  - Midterm analogy: you are cheating on the test.

Golden Rule of Machine Learning

- Even though what we care about is test error:
  - – THE TEST DATA CANNOT INFLUENCE THE TRAINING PHASE IN ANY WAY.

Tom Simonite
June 4, 2015

Why and How Baidu Cheated an Artificial Intelligence Test

Machine learning gets its first cheating scandal.

The sport of training software to act intelligently just got its first cheating scandal. Last month Chinese search company Baidu announced that its image recognition software had inched ahead of Google's on a standardized

http://www.technologyreview.com/view/538111/why-and-how-baidu-cheated-an-artificial-intelligence-test/

Golden Rule of Machine Learning

- Even though what we care about is test error:
  - THE TEST DATA CANNOT INFLUENCE THE TRAINING PHASE IN ANY WAY.

- You also shouldn't change the test set to get the result you want.

DECEPTION AT DUKE: FRAUD IN CANCER CARE?

Were some cancer patients at Duke University given experimental treatments based on fabricated data? Scott Pelley reports.

  - http://blogs.sciencemag.org/pipeline/archives/2015/01/14/the_dukepotti_scandal_from_the_inside

https://www.cbsnews.com/news/deception-at-duke-fraud-in-cancer-care/

Is Learning Possible?

- Does training error say anything about test error?
  - In general, NO: Test data might have nothing to do with training data.
  - E.g., adversary takes training data and flips all labels.

X =

Egg	Milk	Fish
0	0.7	0
0.3	0.7	1
0.3	0	0

y =

Sick?
1
1
0

Xtest =

Egg	Milk	Fish
0	0.7	0
0.3	0.7	1
0.3	0	0

ytest =

Sick?
0
0
1

- In order to learn, we need assumptions:
  - The training and test data need to be related in some way.
  - Most common assumption: independent and identically distributed (IID).

IID Assumption

- Training/test data is independent and identically distributed (IID) if:
  - All objects come from the same distribution (identically distributed).
  - The object are sampled independently (order doesn't matter).

Row 1 comes from same distribution as rows 2-3.

Age	Job?	City	Rating	Income
23	Yes	Van	A	22,000.00
23	Yes	Bur	BBB	21,000.00
22	No	Van	CC	0.00
25	Yes	Sur	AAA	57,000.00

Row 4 does not depend on values in rows 1-3.

- Examples in terms of cards:
  - Pick a card, put it back in the deck, re-shuffle, repeat. → IID
  - Pick a card, put it back in the deck, repeat.
  - Pick a card, don't put it back, re-shuffle, repeat.

Not IID

Testing IID?

IID Assumption and Food Allergy Example

- Is the food allergy data IID?
  - Do all the objects come from the same distribution?
  - Does the order of the objects matter?
- No!
  - Being sick might depend on what you ate yesterday (not independent).
  - Your eating habits might changed over time (not identically distributed).
- What can we do about this?
  - Just ignore that data isn't IID and hope for the best?
  - For each day, maybe add the features from the previous day?
  - Maybe add time as an extra feature?

Learning Theory

- Why does the IID assumption make learning possible?
  - Patterns in training examples are likely to be the same in test examples.
- The IID assumption is rarely true:
  - But it is often a good approximation.
  - There are other possible assumptions.

- Learning theory explores how training error is related to test error.
- We'll look at a simple example, using this notation:
  - E_{train} is the error on training data.
  - E_{test} is the error on testing data.

Fundamental Trade-Off

- Start with $E_{test} = E_{test}$, then add and subtract E_{train} on the right:

$$\underbrace{E_{test}}_{\text{"test error"}} = \underbrace{(E_{test} - E_{train})}_{\text{Difference between training and test error.}} + \underbrace{E_{train}}_{\text{"training error"}}$$

- How does this help?
  - If $(E_{test} - E_{train})$ is not too big, then E_{train} is a good approximation to E_{test}.
- What makes (Etest – Etrain) big?
  - It grows as the model get more "complicated".

Fundamental Trade-Off

- This leads to a fundamental trade-off:
  1. How small you can make the training error.

     vs.

  2. How well training error approximates the test error.

- Simple models (like decision stumps):
  - Training error is a good approximation of test error (not sensitive to training set).
  - But may not fit training data well.
- Complex models (like deep decision trees):
  - Fit training data well.
  - Training error may be poor approximation of test error (sensitive to training set).

Fundamental Trade-Off

- Training error vs. test error for choosing depth:
  - Training error gets better with depth.
  - Test error initially goes down, but eventually increases (overfitting).

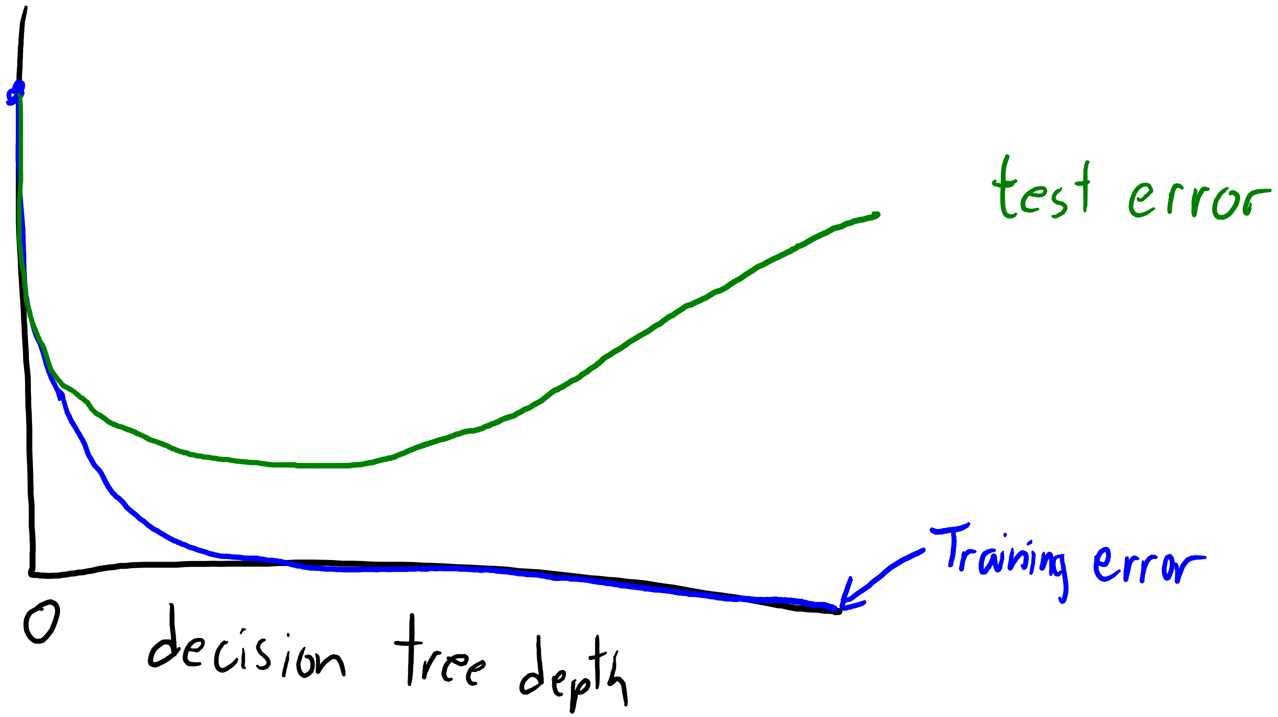

Validation Error

- How do we decide decision tree depth?
- We care about test error.
- But we can't look at test data.
- So what do we do?????

- One answer: Use part of your dataset to approximate test error.
- Split training objects into training set and validation set:
  - Train model based on the training data.
  - Test model based on the validation data.

Validation Error

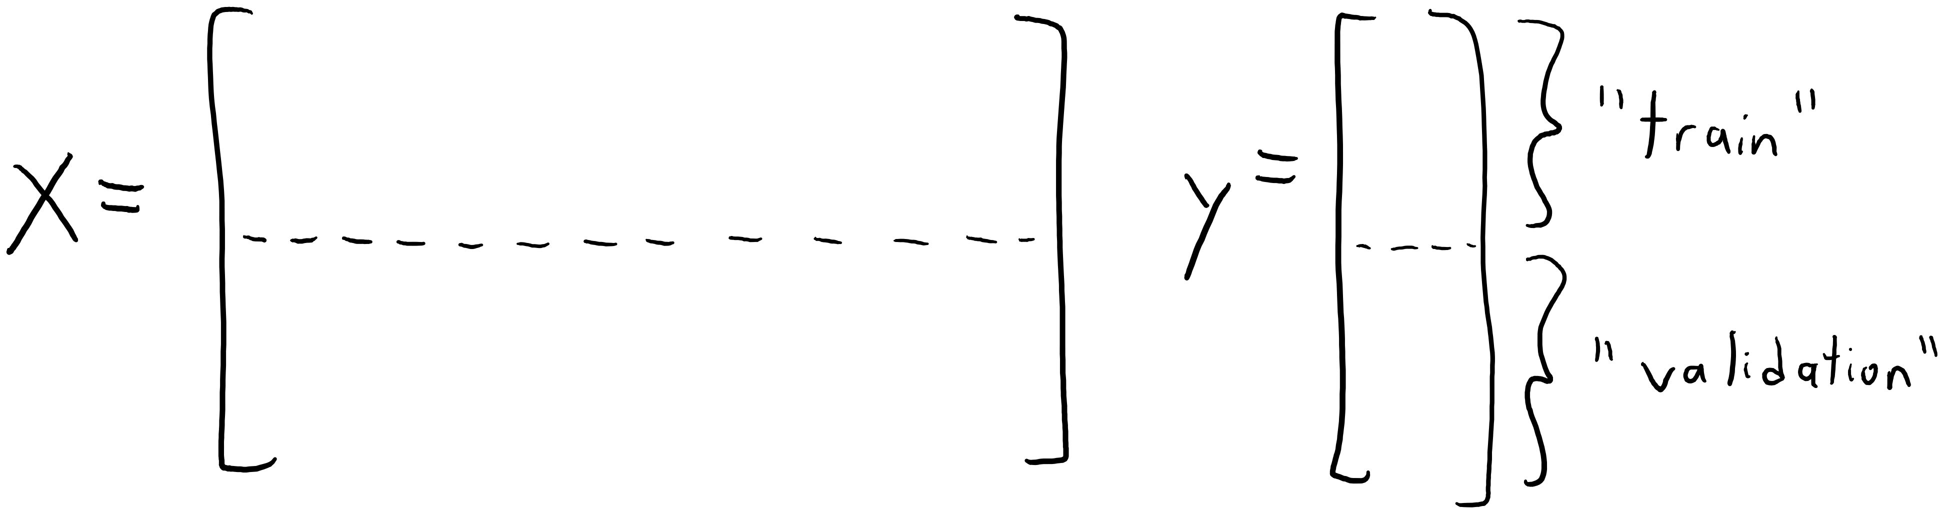

Step 1 is <u>training</u>: $\text{model} = \text{train}(X_{train}, y_{train})$

Step 2 is <u>predicting</u>: $\hat{y} = \text{predict}(\text{model}, X_{validate})$

Step 3 is <u>validating</u>: $\text{error} = \text{sum}(\hat{y} \neq y_{validate})$

Note: if examples are ordered, split should be <u>random</u>.

Validation Error

- Validation error gives an unbiased approximation of test error.

- Midterm analogy:
  - You have 2 practice midterms.
  - You hide one midterm, and spend a lot of time working through the other.
  - You then do the other practice term, as an approximation of how you will do on the test.

- This leads to the following practical strategy:
  - Try a depth-1 decision tree, compute validation error.
  - Try a depth-2 decision tree, compute validation error.
  - Try a depth-3 decision tree, compute validation error.
  - …
  - Choose the depth with the lowest validation error.

Validation Error

- Validation error is much less likely to lead to overfitting.
  - We optimize the validation error over a view values of “depth”.
  - Unlike training error, where we optimize over tons of decision trees.

- But it can still overfit (very common in practice):
  - Validation error is only an unbiased approximation if you use it once.
  - If you minimize it to choose between models, introduces optimization bias:
    - If you try lots of models, one might get a low test error by chance.

Summary

- Training error vs. testing error:
  - What we care about in machine learning is the testing error.
- Golden rule of machine learning:
  - The test data cannot influence training the model in any way.
- Independent and identically distributed (IID):
  - One assumption that makes learning possible.
- Fundamental trade-off:
  - Trade-off between getting low training error and having training error approximate test error.
- Validation set:
  - We can save part of our training data to approximate test error.

- Next time:
  - We discuss the "best" machine learning method.

Bounding (E_{test} – E_{train})

- Let's assume we have a fixed model 'h' (like a decision tree), and then we collect a training set of 'n' examples.
- What is the probability that the error on this training set (E_{train}), is within some small number ε of the test error (E_{test})?

- From "Hoeffding's inequality" we have:

$$P[|E_{train}(h) - E_{test}(h)| > \varepsilon] \leq 2\exp(-2\varepsilon^2 n)$$

- This is great! In this setting the probability that our training error is far from our test error goes down exponentially in terms of the number of samples 'n'.

Bounding ($E_{test} – E_{train}$)

- Unfortunately, the last slide gets it backwards:
  - We usually don’t pick a model and then collect a dataset.
  - We usually collect a dataset and then pick the model ‘w’ based on the data.
- We now picked the model that did best on the data, and Hoeffding’s inequality doesn’t account for the optimization bias of this procedure.

- One way to get around this is to bound ($E_{test} – E_{train}$) for *all* models in the space of models we are optimizing over.
  - If bound it for all models, then we bound it for the best model.
  - This gives looser but correct bounds.

Bounding (E_{test} – E_{train})

- If we only optimize over a finite number of events ‘k’, we can use the “union bound” that for events $\{A_1, A_2, \ldots, A_k\}$ we have:

$$p(A_1 \cup A_2 \cup \cdots \cup A_k) \leq \sum_{i=1}^{k} p(A_i)$$

- Combining Hoeffding’s inequality and the union bound gives:

$$p\left(|E_{train}(h_1) - E_{test}(h_1)| > \varepsilon \cup |E_{train}(h_2) - E_{test}(h_2)| > \varepsilon \cup \cdots \cup |E_{train}(h_k) - E_{test}(h_k)| > \varepsilon\right)$$

$$\leq \sum_{i=1}^{k} p\left(|E_{train}(h_i) - E_{test}[h_i]| > \varepsilon\right)$$

$$\leq \sum_{i=1}^{k} 2\exp(-2\varepsilon^2 n)$$

$$\leq 2k\exp(-2\varepsilon^2 n)$$

Bounding (E_{test} – E_{train})

- So, with the optimization bias of setting "h*" to the best 'h' among 'k' models, probability that (Etest – Etrain) is bigger than ε satisfies:

$$p(|E_{train}(h^*) - E_{test}(h^*)| > \epsilon) \leq 2k\exp(-2\epsilon^2 n)$$

- So optimizing over a few models is ok if we have lots of examples.
- If we try lots of models then (E_{test} – E_{train}) could be very large.
- Later in the course we'll be searching over continuous models where k = infinity, so this bound is useless.
- To handle continuous models, one way is via the VC-dimension.
  - Simpler models will have lower VC-dimension.

Refined Fundamental Trade-Off

- Let E_{best} be the irreducible error (lowest possible error for *any* model).
  - For example, irreducible error for predicting coin flips is 0.5.
- Some learning theory results use E_{best} to futher decompose E_{test}:

$$E_{test} = \underbrace{(E_{test} - E_{train})}_{\text{"variance"}} + \underbrace{(E_{train} - E_{best})}_{\text{"bias"}} + \underbrace{E_{best}}_{\text{"noise"}}$$

- This is similar to the bias-variance decomposition:
  - Term 1: measure of variance (how sensitive we are to training data).
  - Term 2: measure of bias (how low can we make the training error).
  - Term 3: measure of noise (how low can any model make test error).

Refined Fundamental Trade-Off

- Decision tree with high depth:
  - Very likely to fit data well, so bias is low.
  - But model changes a lot if you change the data, so variance is high.
- Decision tree with low depth:
  - Less likely to fit data well, so bias is high.
  - But model doesn’t change much you change data, so variance is low.

- And degree does not affect irreducible error.
  - Irreducible error comes from the best possible model.

Bias-Variance Decomposition

- Analysis of expected test error of any learning algorithm:

Assume $y_i = f(x_i) + \epsilon$, for some function 'f' and random error ϵ with a mean of 0 and a variance of σ^2.

Assume we have a "learner" that can take a training set $\{(x_1, y_1), (x_2, y_2), \ldots, (x_n, y_n)\}$ and use these to make predictions $\hat{f}(x_i)$.

Then for a new example (x_i, y_i) the error averaged over training sets is

$$E[(y_i - \hat{f}(x_i))^2] = Bias[\hat{f}(x_i)]^2 + Var[\hat{f}(x_i)] + \sigma^2$$

"Irreducible error": best we can hope for given the noise level.

Where $Bias[\hat{f}(x_i)] = E[\hat{f}(x_i)] - f(x_i)$,

Expected error due to having wrong model.

$Var[\hat{f}(x_i)] = E[(\hat{f}(x_i) - E[\hat{f}(x_i)])^2]$

How sensitive is the model to the particular training set?

Learning Theory

- Bias-variance decomposition is a bit weird compared to our previous decompositions of E_{test}:
  - Bias-variance decomposition considers expectation over *possible training sets*.
  - But doesn't say anything about test error with *your* training set.

- Some keywords if you want to learn about learning theory:
  - Bias-variance decomposition, sample complexity, probably approximately correct (PAC) learning, Vapnik-Chernovenkis (VC) dimension, Rademacher complexity.

- A gentle place to start is the "Learning from Data" book:
  - https://work.caltech.edu/telecourse.html

A Theoretical Answer to "How Much Data?"

- Assume we have a source of IID examples and a fixed class of parametric models.
  - Like "all depth-5 decision trees".
- Under some nasty assumptions, with 'n' training examples it holds that:

  E[test error of best model on training set] – (best test error in class) = O(1/n).

- You rarely know the constant factor, but this gives some guidelines:
  - Adding more data helps more on small datasets than on large datasets.
    - Going from 10 training examples to 20, difference with best possible error gets cut in half.
      - If the best possible error is 15% you might go from 20% to 17.5% (this does **not** mean 20% to 10%).
    - Going from 110 training examples to 120, error only goes down by ~10%.
    - Going from 1M training examples to 1M+10, you won't notice a change.
  - Doubling the data size cuts the error in half:
    - Going from 1M training to 2M training examples, error gets cut in half.
    - If you double the data size and your test error doesn't improve, more data might not help.